

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: df= pd.read_csv("C:\\Users\\DELL\\Downloads\\archive (1)\\diamonds.csv")
```

```
In [3]: df.head()
```

Out[3]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75

```
In [4]: df.tail()
```

Out[4]:

	carat	cut	color	clarity	depth	table	price	x	y	z
53935	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50
53936	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61
53937	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56
53938	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53940 entries, 0 to 53939
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   carat       53940 non-null  float64
1   cut         53940 non-null  object
2   color       53940 non-null  object
3   clarity     53940 non-null  object
4   depth       53940 non-null  float64
5   table       53940 non-null  float64
6   price       53940 non-null  int64
7   x           53940 non-null  float64
8   y           53940 non-null  float64
9   z           53940 non-null  float64
dtypes: float64(6), int64(1), object(3)
memory usage: 4.1+ MB
```

In [6]: `df.describe()`

Out[6]:

	carat	depth	table	price	x	y
count	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000	539
mean	0.797940	61.749405	57.457184	3932.799722	5.731157	5.734526
std	0.474011	1.432621	2.234491	3989.439738	1.121761	1.142135
min	0.200000	43.000000	43.000000	326.000000	0.000000	0.000000
25%	0.400000	61.000000	56.000000	950.000000	4.710000	4.720000
50%	0.700000	61.800000	57.000000	2401.000000	5.700000	5.710000
75%	1.040000	62.500000	59.000000	5324.250000	6.540000	6.540000
max	5.010000	79.000000	95.000000	18823.000000	10.740000	58.900000

In [7]: `df.shape`

Out[7]: (53940, 10)

In [8]: `df.isnull().sum()`

Out[8]:

```
carat      0
cut        0
color      0
clarity    0
depth      0
table      0
price      0
x          0
y          0
z          0
dtype: int64
```

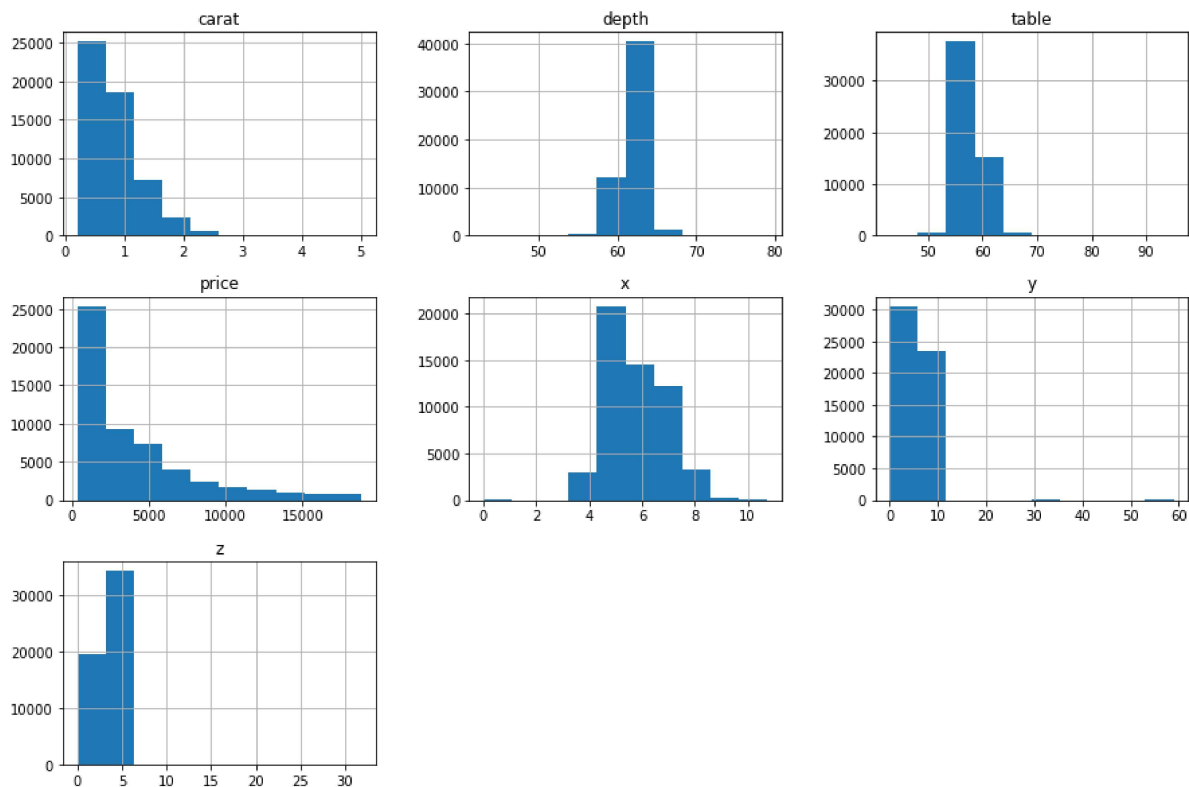
EDA

In [9]: `num_features = df.select_dtypes(include=['int64','float64'])`
`print(num_features.columns)`

Index(['carat', 'depth', 'table', 'price', 'x', 'y', 'z'], dtype='object')

```
In [10]: num_features.hist(figsize =(15,10))
```

```
Out[10]: array([[<AxesSubplot:title={'center':'carat'}>,
  <AxesSubplot:title={'center':'depth'}>,
  <AxesSubplot:title={'center':'table'}>],
  [<AxesSubplot:title={'center':'price'}>,
  <AxesSubplot:title={'center':'x'}>,
  <AxesSubplot:title={'center':'y'}>],
  [<AxesSubplot:title={'center':'z'}>, <AxesSubplot:>,
  <AxesSubplot:>]], dtype=object)
```




```
In [11]: cat_features = df.select_dtypes(include=['object'])
print(cat_features.columns)
```

```
Index(['cut', 'color', 'clarity'], dtype='object')
```


```
In [12]: plt.figure(figsize=(24, 48))

plt.subplot(411)
sns.countplot(y='cut', hue='color', data = cat_features)
plt.title('Color vs Cut')
```

Out[12]: Text(0.5, 1.0, 'Color vs Cut')




```
In [13]: for num_var in num_features:
plt.figure(figsize=(10,5))
sns.boxplot(x=df[num_var])
```





```
In [14]: corr=df.corr()
sns.heatmap(corr,annot=True)
```

Out[14]: <AxesSubplot:>



```
In [15]: for num_var in num_features:
plt.figure(figsize=(5,5))
plt.xlabel('price')
plt.ylabel(num_var)
x=df['price']
y=df[num_var]
plt.scatter(x,y)
```



Feature Transformation

```
In [16]: df_cat = pd.get_dummies(cat_features, drop_first=True)
df_cat.head()
```

Out[16]:

	cut_Good	cut_Ideal	cut_Premium	cut_Very Good	color_E	color_F	color_G	color_H	color_I	color_J
0	0	1	0	0	1	0	0	0	0	0
1	0	0	1	0	1	0	0	0	0	0
2	1	0	0	0	1	0	0	0	0	0
3	0	0	1	0	0	0	0	0	0	1
4	1	0	0	0	0	0	0	0	0	0

```
In [17]: new_df = pd.concat([num_features, df_cat], axis=1)
new_df.head()
```

Out[17]:

	carat	depth	table	price	x	y	z	cut_Good	cut_Ideal	cut_Premium	...	color_H	color_I
0	0.23	61.5	55.0	326	3.95	3.98	2.43	0	1	0	...	0	0
1	0.21	59.8	61.0	326	3.89	3.84	2.31	0	0	1	...	0	0
2	0.23	56.9	65.0	327	4.05	4.07	2.31	1	0	0	...	0	0
3	0.29	62.4	58.0	334	4.20	4.23	2.63	0	0	1	...	0	0
4	0.31	63.3	58.0	335	4.34	4.35	2.75	1	0	0	...	0	0

5 rows × 24 columns

```
In [18]: X = new_df.drop(columns = ['price'], axis = 1)
X.head()
```

Out[18]:

	carat	depth	table	x	y	z	cut_Good	cut_Ideal	cut_Premium	cut_Very Good	...	color_H	color_I
0	0.23	61.5	55.0	3.95	3.98	2.43	0	1	0	0	...	0	0
1	0.21	59.8	61.0	3.89	3.84	2.31	0	0	1	0	...	0	0
2	0.23	56.9	65.0	4.05	4.07	2.31	1	0	0	0	...	0	0
3	0.29	62.4	58.0	4.20	4.23	2.63	0	0	1	0	...	0	0
4	0.31	63.3	58.0	4.34	4.35	2.75	1	0	0	0	...	0	0

5 rows × 23 columns

```
In [19]: y = new_df['price']
y.head()
```

```
Out[19]: 0    326
         1    326
         2    327
         3    334
         4    335
         Name: price, dtype: int64
```

Splitting

```
In [20]: # split into train and test
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=42)
```

```
In [21]: print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

(37758, 23) (37758,)
(16182, 23) (16182,)
```

```
In [22]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```
In [23]: scaler.fit_transform(X_train)
scaler.transform(X_test)
```

```
Out[23]: array([[0.07692308, 0.51388889, 0.26923077, ..., 1.          , 0.          ,
                0.          ],
               [0.1995842 , 0.51388889, 0.23076923, ..., 0.          , 0.          ,
                0.          ],
               [0.06444906, 0.56111111, 0.28846154, ..., 0.          , 0.          ,
                0.          ],
               ...,
               [0.08523909, 0.50833333, 0.23076923, ..., 0.          , 0.          ,
                0.          ],
               [0.04365904, 0.525      , 0.25      , ..., 0.          , 0.          ,
                0.          ],
               [0.16216216, 0.51666667, 0.44230769, ..., 0.          , 0.          ,
                0.          ]])
```

Training the Model

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
linear_model = LinearRegression()  
linear_model.fit(X_train, y_train)
```

Out[24]: LinearRegression()

```
In [25]: y_test_pred = linear_model.predict(X_test)
```

```
In [26]: from sklearn import metrics
```

```
In [27]: metrics.r2_score(y_test, y_test_pred)
```

Out[27]: 0.9182661537265939

Random Forest Regressor

```
In [28]: from sklearn.ensemble import RandomForestRegressor
```

```
In [29]: model_rf = RandomForestRegressor()
```

```
In [30]: model_rf.fit(X_train, y_train)
```

Out[30]: RandomForestRegressor()

```
In [31]: y_pred_rf = model_rf.predict(X_test)
```

```
In [32]: metrics.r2_score(y_test, y_pred_rf)
```

Out[32]: 0.9739244252159122

Decision Tree

```
In [33]: from sklearn.tree import DecisionTreeRegressor
```

```
In [34]: model_tree = DecisionTreeRegressor()
```

```
In [35]: model_tree = DecisionTreeRegressor()
```

```
In [36]: model_tree.fit(X_train, y_train)
```

Out[36]: DecisionTreeRegressor()

```
In [37]: y_pred_tree = model_tree.predict(X_test)
```


In [38]: `metrics.r2_score(y_test,y_pred_tree)`

Out[38]: 0.9524367348772984

In []: